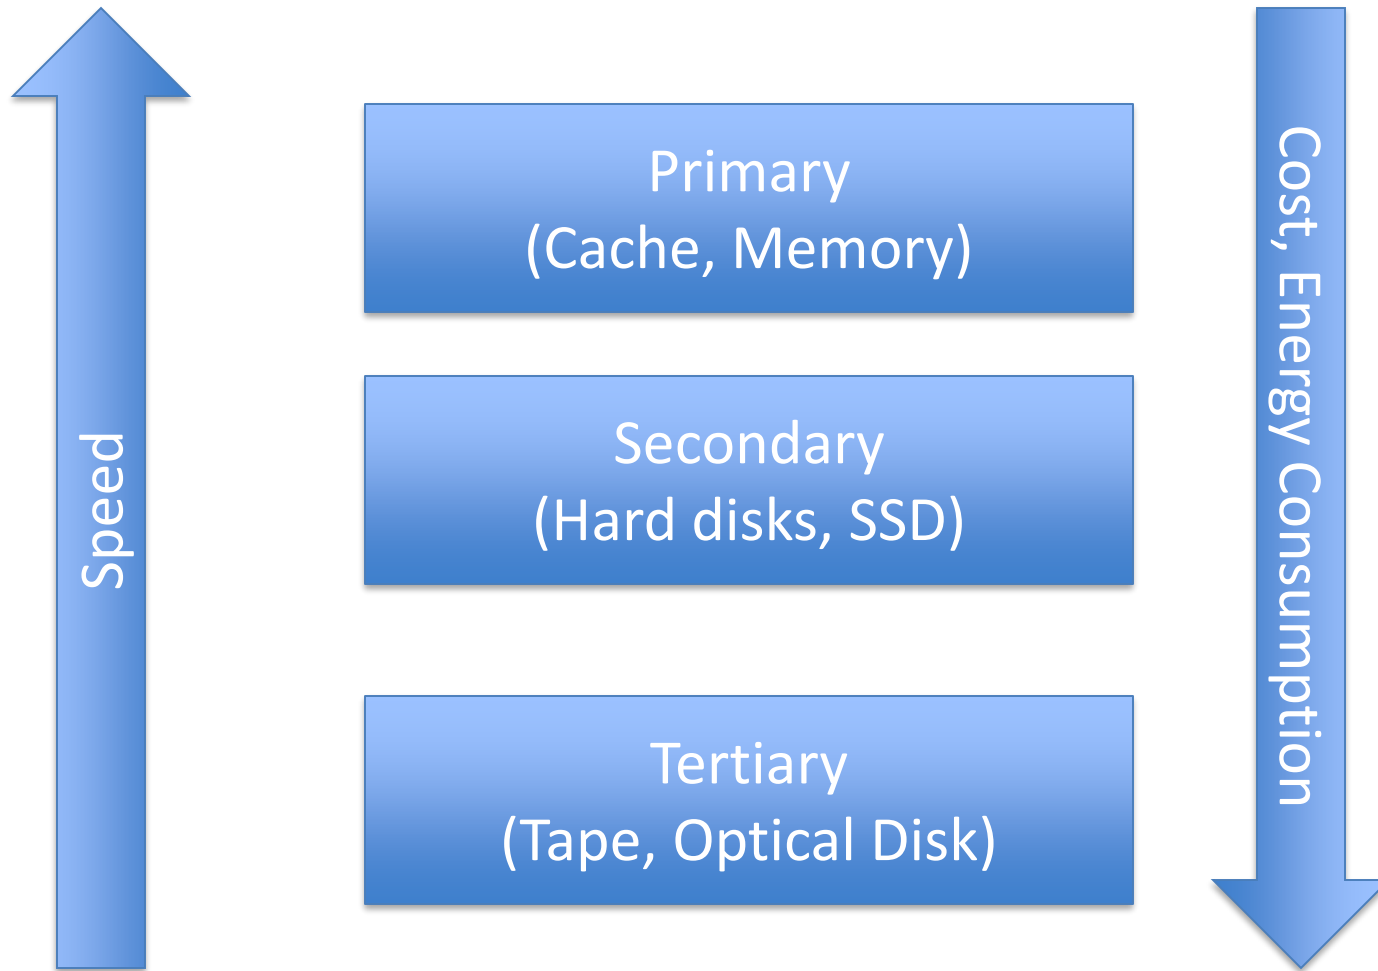


Storage Systems

DSCI 551

Wensheng Wu

Storage hierarchy




Access times



Time taken before drive is ready to transfer data

LEVEL	ACCESS TIME	TYPICAL SIZE
Registers	"instantaneous"	under 1KB
Level 1 Cache	1-3 ns	64KB per core
Level 2 Cache	3-10 ns	256KB per core
Level 3 Cache	10-20 ns	2-20 MB per chip
Main Memory	30-60 ns	4-32 GB per system
Hard Disk	3,000,000-10,000,000 ns	over 1TB



SSD: 25,000 ns

Resource: <https://arstechnica.com/information-technology/2012/06/inside-the-ssd-revolution-how-solid-state-disks-really-work/>

CRUD

- Basic functions of a storage device
- CRUD:
 - (C)reate/write
 - (R)ead
 - (U)pdate/overwrite
 - (D)elele

Characterizing a storage device

- Capacity (bytes)
 - How much data it can hold
- Cost (\$\$\$)
 - Price per byte of storage
- Bandwidth (bytes/sec)
 - Number of bytes that can be transferred per second
 - Note that read and write bandwidth may be different
- Latency (seconds)
 - Time elapsed, waiting for response/delivery of data

Time to complete an operation


- Time to complete an operation depends on both bandwidth and latency
 - $\text{CompletionTime} = \text{Latency} + \text{Size}/\text{Bandwidth}$
- The time for a workload may depend on
 - Technology, e.g., hard drive/SSD
 - Operation type, e.g., read/write
 - Number of operations in the workload
 - Access pattern (random vs. sequential)

Access pattern

- Sequential
 - Data to be accessed are located next to each other or sequentially on the device
- Random
 - Access data located randomly on storage device


Magnetic recording

- Write head
 - Applies electrical current to write head
 - Changes direction of magnetic field under head



Reading

- Read head senses direction of magnetic field



Road map


- Tapes



- Hard disk
- Solid state drive

Linear tape

- Data recorded on parallel tracks that span the length of the tape



Tapes

- Current technology is LTO
 - Linear Tape-Open (an open standard)
- Characteristics
 - Capacity up to 6.25TB per tape (LTO-7)
 - Drive cost ~ \$2000
 - Tape cost ~ \$60 for 6TB tape
- Tape access time (~ minute)
 - Time to mount the tape
 - Time to wind the tape to correct position
- Data transmission rates ~ 250MB/sec



\$60
300MB/s
6TB
LTO-7

Performance characteristics

- High latency/low cost makes tape most appropriate for "archival" storage
 - Low frequency of (mostly sequential) reads
 - Very large data objects
- Random access will be slow due to latency
 - Sequential reads will be fast

Linear tape file system

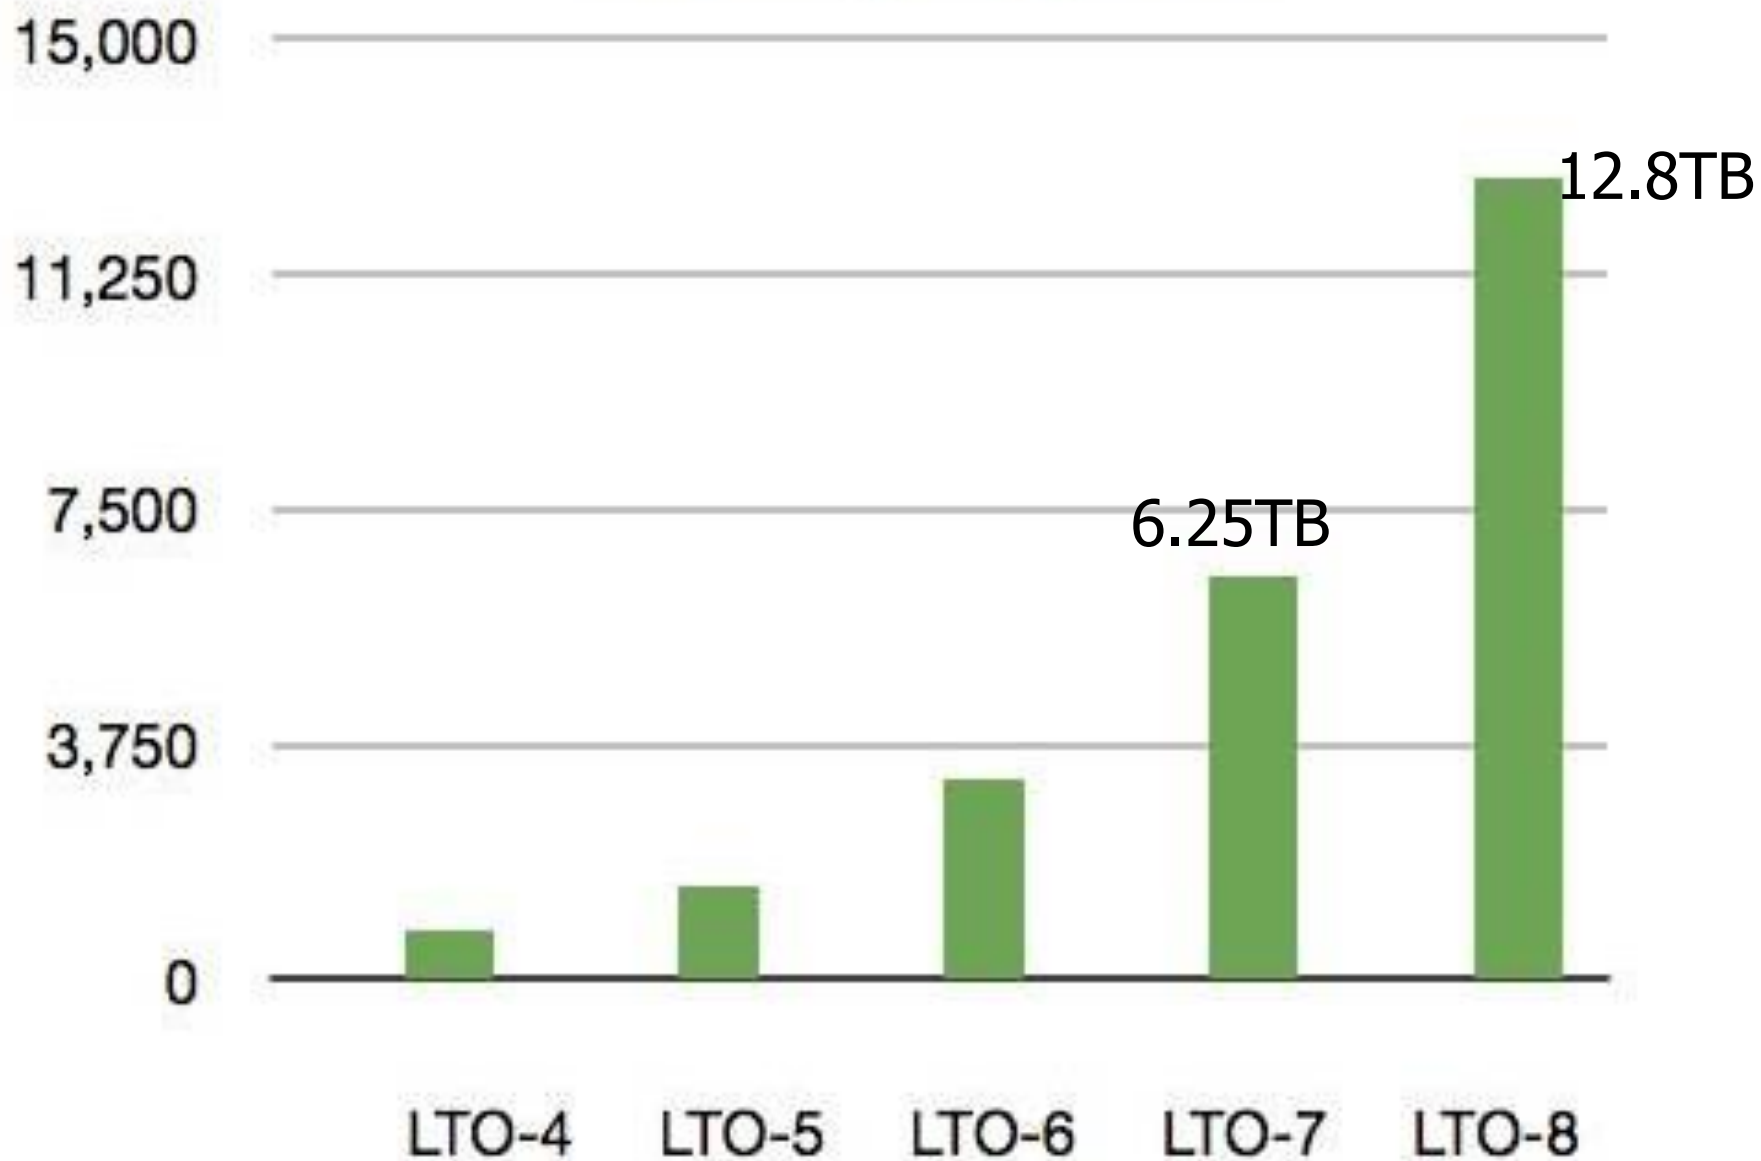
- Two partitions on tape
 - First contains metadata and directories. Tape reader can find and load this very quickly
 - Second contains blocks for data
- Directory structure coded in XML
 - Self describing file format...



Tape Cartridge



Raw Capacity in GB



A tape library




Inside a robotic tape library

- <https://www.youtube.com/watch?v=nYfTtvpQ778>



Road map

- Tapes
- Hard disk 
- Solid state drive

Hard disk drives

- Perhaps the most pervasive form of storage
- Basic Idea:
 - One or more spinning magnetic platters
 - Typically two surfaces per platter
 - Disk arm positions over the radial position (tracks) where data are stored
 - It swings across tracks (but do not extend/shrink)
 - Data is read/written by a read/write head as platter spins



Sponsored ⓘ

WD 4TB 3.5 Inch SATA III, 7200 RPM, 64 MB Cache Enterprise Hard Drive (WD4000FYYZ)

★★★★☆ ~ 321

\$123⁰⁰

✓prime FREE Delivery Thu, Jan 23

Amazon's Choice



WD Blue 1TB PC Hard Drive - 7200 RPM Class, SATA 6 Gb/s, 64 MB Cache, 3.5" - WD10EZEX

★★★★☆ ~ 24,141

\$44⁰⁹ ~~\$109.99~~

✓prime FREE Delivery Thu, Jan 23

More Buying Choices

\$34.98 (76 used & new offers)

Internal of hard disk

Actuator

Spindle




Platter

Disk head

Disk arm and platter




Disk arm & head close-up



Disk head close-ups





Disk head movement

- Hard disk head movement while copying files between two folders (e.g., partition c to d)
 - <https://www.youtube.com/watch?v=BlB49F6ExkQ>



2GB Storage in 1980s (\$250,000!)



Physical characteristics


- 3.5" (diameter, common in desktops)
- 2.5" (common in laptops)
- Rotational speed
 - 4,800 RPM
 - 5,400 RPM
 - 7,200 RPM
 - 10,000 RPM
- Between 5-7 platters
- Current capacity up to 10TB (Western Digital)

Disk organization

- Each platter consists of a number of tracks
- Each track is divided into N fixed size sectors
 - Typical sector size is 512 bytes (old) or 4KB (new)
 - Sectors can be numbered from 0 to N-1
 - Entire sector is written "atomically"
 - All or nothing


CHS (cylinder-head-sector)

- Early way to address a sector
 - Now LBA ([Logical Block Addressing](https://en.wikipedia.org/wiki/Logical_Block_Addressing)) more common



<https://en.wikipedia.org/wiki/Cylinder-head-sector>

CHS (Wikipedia)




Example

- # cylinders: 256
- # heads: 16 (i.e., 8 platters, 2 heads/platter)
- # sectors/track: 64
- Sector size = 4KB


What is the capacity of the drive?

A simple disk drive (one track only)



Rotational latency

- Waiting for the right sector to rotate under the head
 - On average: about $\frac{1}{2}$ of time of a full rotation
 - Worst case?
 - Best case?




Rotation time

- Assume 10,000 RPM (rotations per minute)

$$\frac{60000 \text{ ms}}{1000 \text{ rotations}} = \frac{6 \text{ ms}}{\text{rotation}}$$

Multiple tracks: add seek times



Average seek time is about 1/3 max seek time
(see reading: Chapter 37, page 9 for more details)

Transfer time

- Assume that we transfer 512KB
- Assume 128 MB/sec transmission bandwidth
- Transfer time:
$$512\text{KB}/128\text{MB} * 1000\text{ms} = 4\text{ms}$$

Completion time

- $T = T_{\text{seek}} + T_{\text{rotation}} + T_{\text{transfer}}$
 - T_{seek} : Time to get the disk head on right track
 - T_{rotation} : Time to wait for the right sector to rotate under the head
 - T_{transfer} : Time to actually transfer the data

Example

- Capacity 4TB
- # platters: 4
- # heads: 8
- Bytes per sector: 4096
- Transmission bandwidth: 100MB/sec
- Maximum seek time: 12ms
- RPM: 10,000

Time to transfer a file

- The file occupies 100 sectors (sequentially)
- Avg. seek time =?
- Avg. rotational latency =?
- Transfer time = ?

Sector vs. block

- Block has 1 or more sectors
- Disk typically transfers one block at a time
- We will assume one block = one sector
 - Unless stated otherwise

Sequential operations

- May assume all sectors involved are on same track
 - We may need to seek to the right track or rotate to the first sector
- But no rotation/seeking needed afterward

Actual bandwidth

- Consider a workload w
 - E.g., w = sequential access of 100 blocks of data
 - Denote size (# of bytes) of data in w as $|w|$
 - E.g., w = 400KB (100 blocks, 4KB/block)
- Suppose completion time for w = t
- Actual bandwidth (with respect to w) = $|w|/t$

Sequential vs. random

- Consider disk with 7ms avg seek, 10,000 RPM platter speed and 50 MB/sec transfer rate, 4KB/block
- Sequential access of 10 MB
 - Completion time = $7 + 3 + 10/50 * 1000 = 210\text{ms}$
 - Actual bandwidth = $10\text{MB}/210\text{ms} = 47.62 \text{ MB/s}$
- Random access of 10 MB (2,500 blocks)
 - Completion time = $2500 * (7 + 3 + 4/50) = 25.2\text{s}$
 - Actual bandwidth = $10\text{MB} / 25.2\text{s} = .397 \text{ MB/s}$

Road map

- Tapes
- Hard disk
- Solid state drive



Solid State Drive



Memory chips



Solid State Drives

- All electronic, made from flash memory
- Lower energy consumption than hard drive
- More expensive, less capacity
 - 3 times or more expensive
- Limited lifetime, can only write a limited number of times.
 - E.g., 100K program/write cycles for SLC (single-level cell) memory

SSD vs. Hard Drive (price)

Amazon's **Choice**



Samsung SSD 860 EVO 1TB 2.5 Inch SATA III Internal SSD (MZ-76E1T0B/AM)

★★★★★ ~ 13,615

\$146⁴³ ~~\$199.99~~

✓prime FREE One-Day

Get it **Tomorrow, Jan 20**

More Buying Choices

\$123.09 (37 used & new offers)

Amazon's **Choice**



WD Blue 1TB PC Hard Drive - 7200 RPM Class, SATA 6 Gb/s, 64 MB Cache, 3.5" - WD10EZEX

★★★★☆ ~ 24,141

Personal Computers

\$44⁰⁹ ~~\$109.99~~

✓prime FREE Delivery **Thu, Jan 23**

More Buying Choices

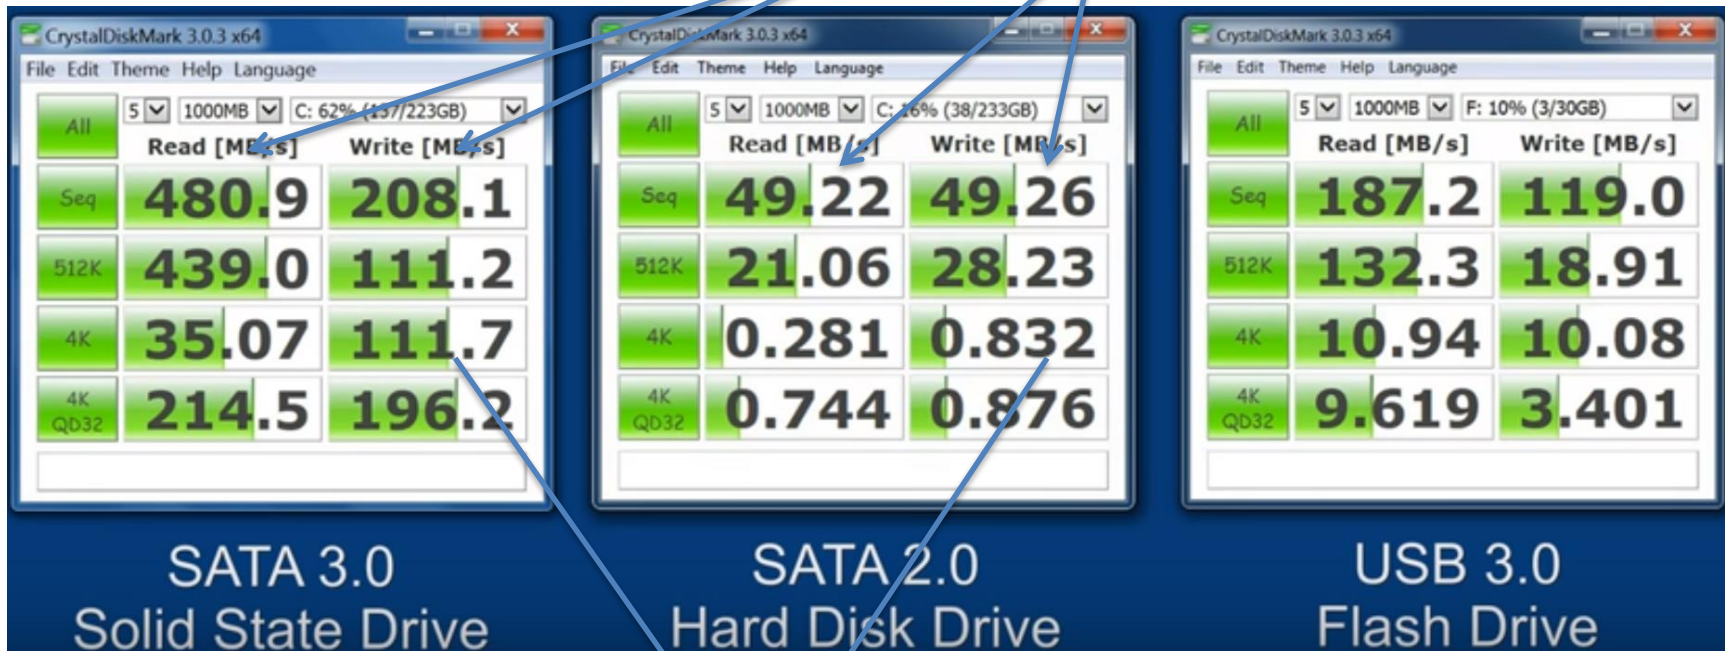
\$34.98 (76 used & new offers)

Solid State Drives

- Same form-factor and control interface as magnetic disks
- Significantly better latency
 - No seek or rotational delay
- Much better performance on random workload:
 - Benefits from improved latency
 - However, write has much higher latency (but see next slide)

Speed comparison ([YouTube](#))

Read vs write: significant difference in SSD vs marginal in HDD



Due to buffered writes


Writing to SSD is complicated

- Can not overwrite a page
 - Need to erase its block (at a certain point) instead
- SSD controllers take care of all these details

SSD

- Contains a number of flash memory chips
 - Chip -> dies -> planes -> blocks -> pages (rows) -> cells
 - Cells are made of floating-gate transistors
- Page is the smallest unit of data transfer between SSD and main memory
 - Much like a block in hard disk

Die Layout



Dies, planes, block, and pages

- Typically, a chip may have 1, 2, or 4 dies
- A die may have 1 or 2 planes
- A plane has a number of blocks
 - Block is the smallest unit that can be erased
- A block has a number of pages
 - Page is the smallest unit that can be read, programmed/written


Typical page and block sizes

- Common page sizes: 2K, 4K, 8K, and 16K
- A block typically has 128 to 256 pages

=> Block size: 256KB to 4MB


Normal transistor (MOSFET)

- When current applied to gate terminal
 - semi-conducting region (purple) becomes conductive




Floating gate transistor

- Contain an additional gate: floating gate
 - Floating, since isolated by oxide layer
 - (thus not connected to other components)



Floating gate transistor

- By applying high positive/negative voltage to control gate, electrons can be attracted to or repelled from floating gate



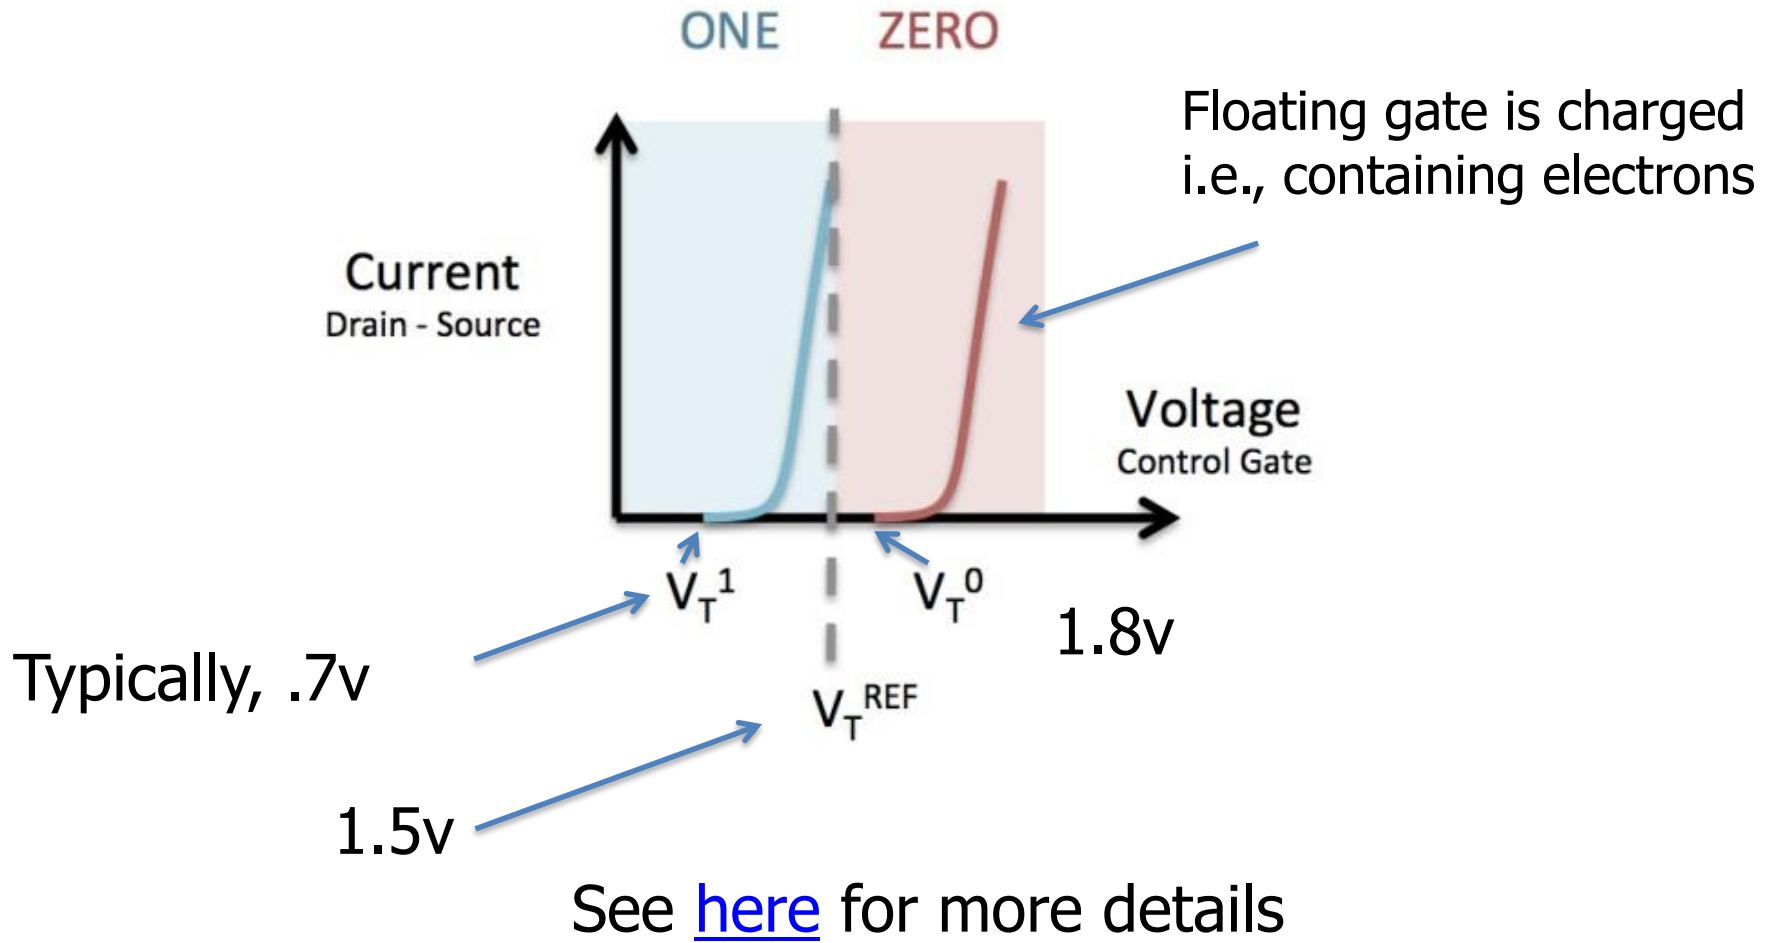
Floating gate transistor

- State = 1, if no electrons in the floating gate
- State = 0, if there are electrons (negative charges)
 - Electrons stuck there even when power is off
 - So state is retained

Read operations

- Electrons on the floating gate affect the threshold voltage for the floating gate transistor to conduct
- Higher voltage needed when gate has electrons

Read operations

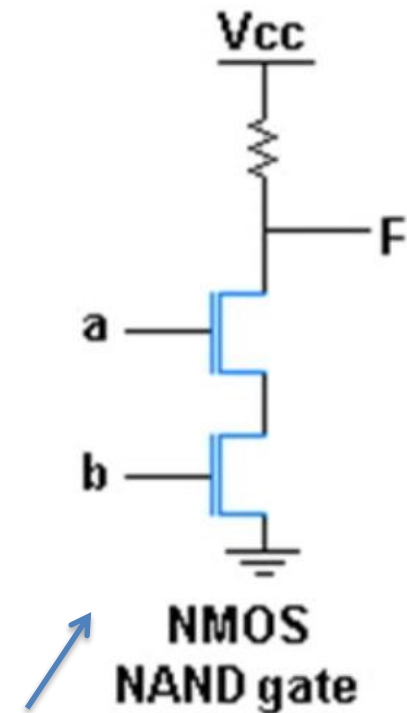
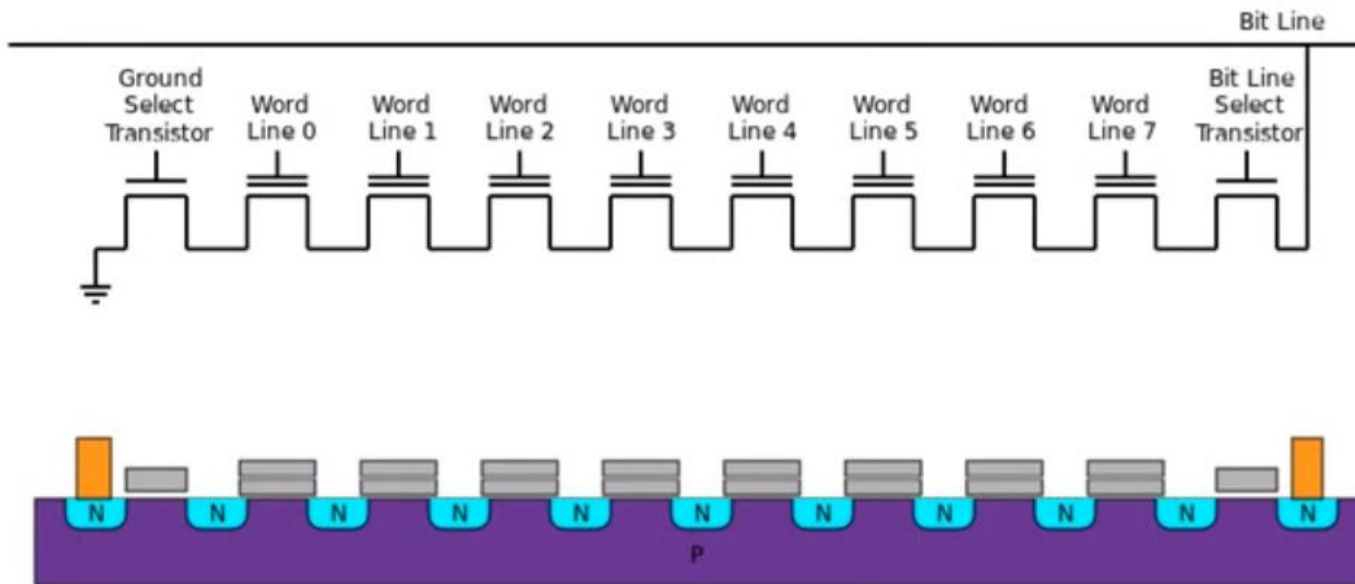


Read operations

- Apply V_{int} (intermediate voltage)
- If the current is detected, gate has no electrons
=> bit = 1
- If no current, gate must have electrons
=> bit = 0


NAND flash layout

- Transistors are strung together in a series
 - Similar to the transistors in an NAND gate

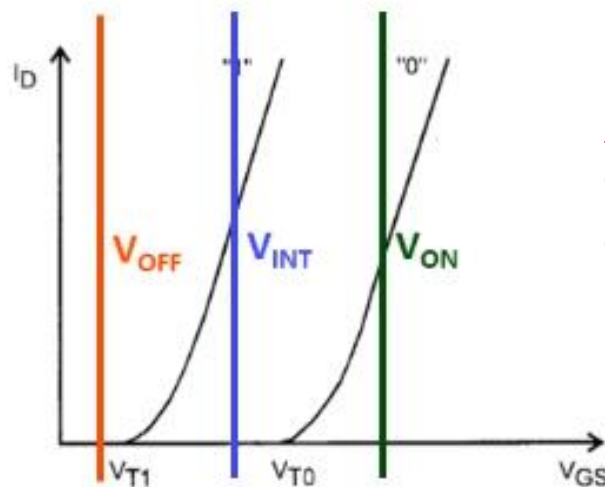


Output $F = 0$ only when both a & $b = 1$
I.e., F will be grounded only when both a and b conduct

NAND reading

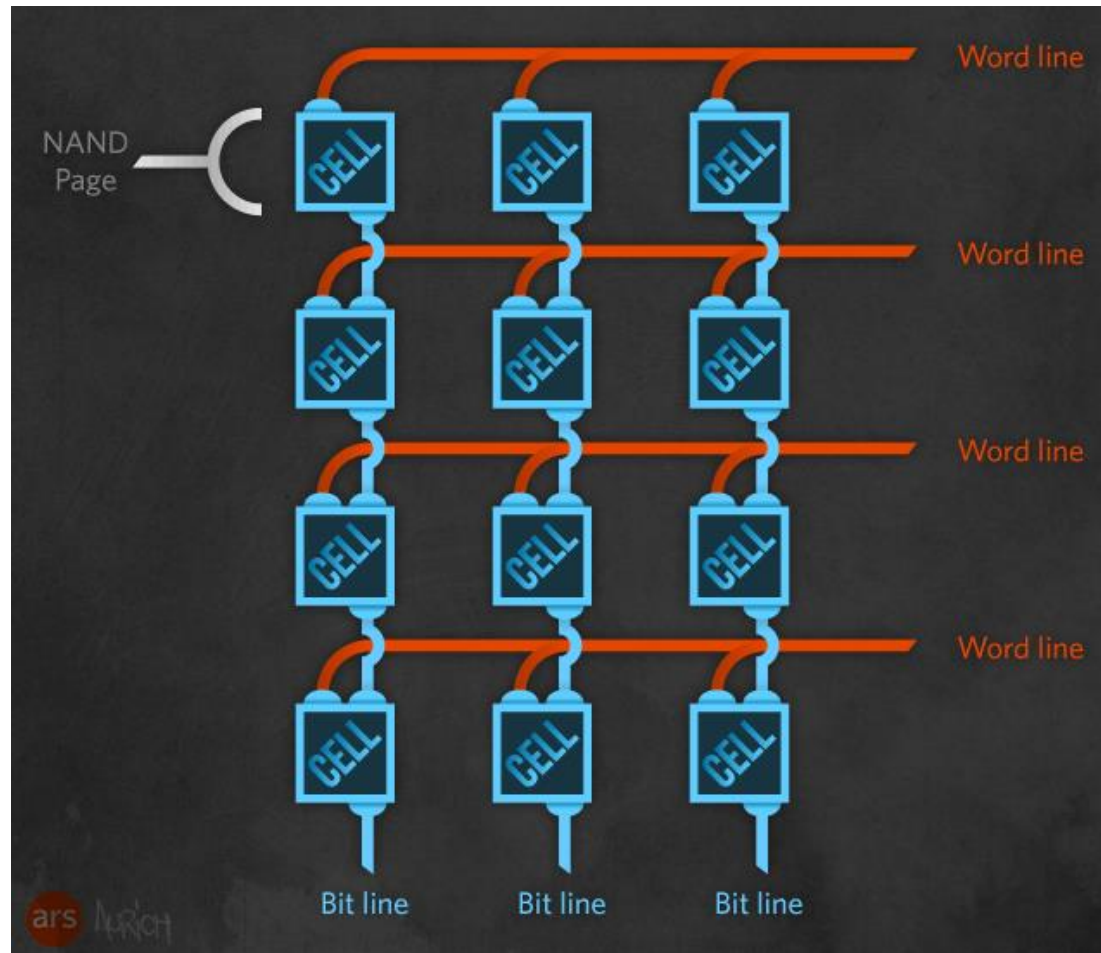


Bit line




Apply V_{on} to all others so that they all conduct, no matter they are charged or not.

NAND flash



Write and erase



- Write: $1 \Rightarrow 0$
 - Apply high POSITIVE voltage (\gg voltage for read) to the control gate
 - Attract electrons from channel to floating gate (through quantum tunneling)
- Erase: $0 \Rightarrow 1$
 - Need to apply much higher NEGATIVE voltage
 - Get rid of electrons from floating gate
 - May stress surrounding cells
 - So dangerous to do on individual pages

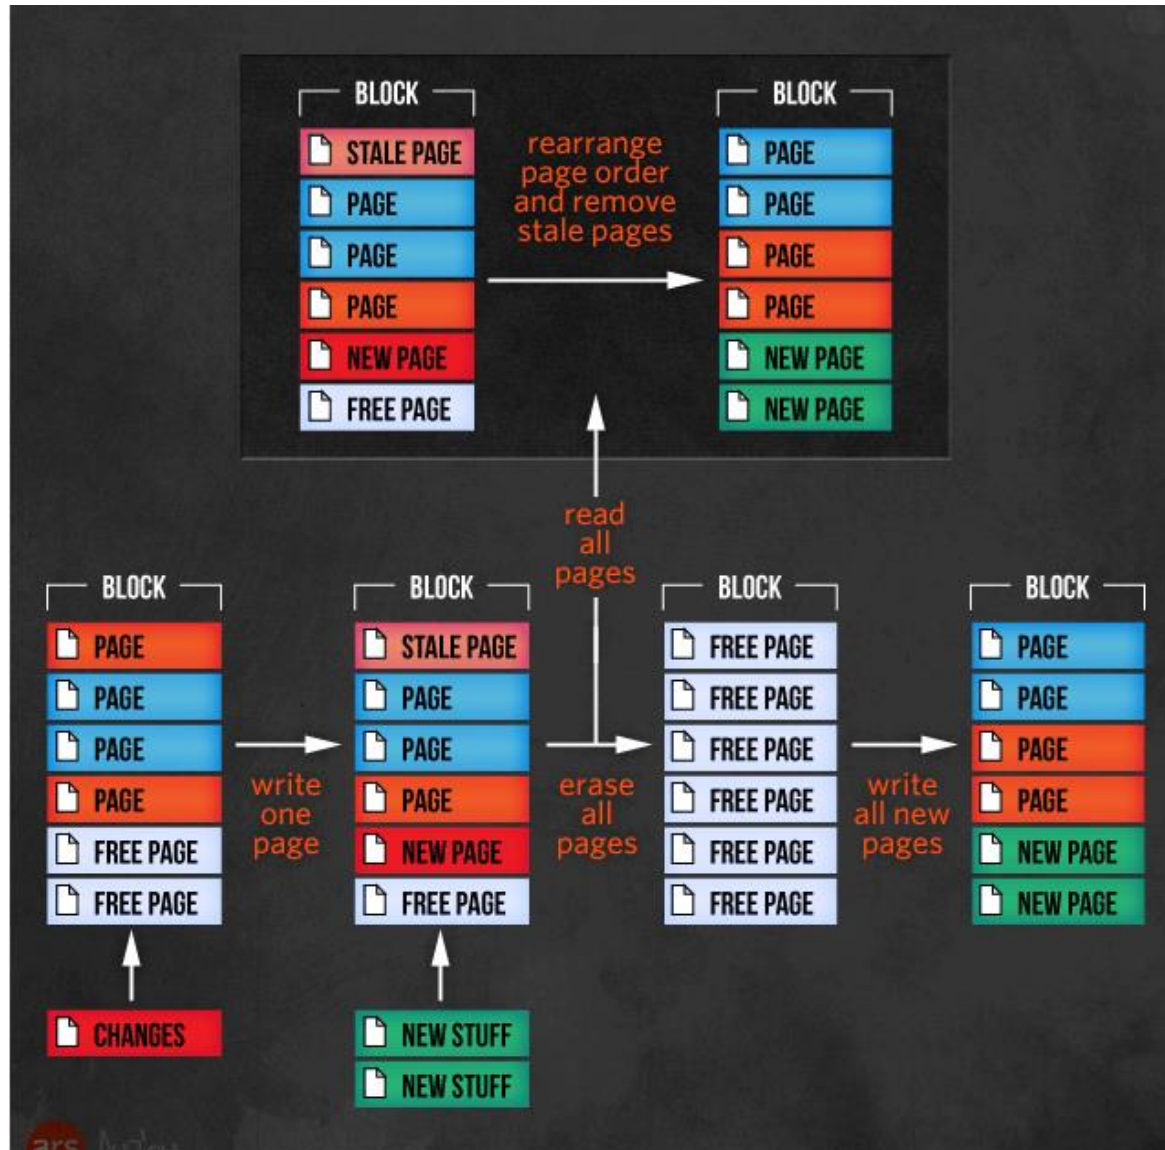
Read/write units

- **Page** is the smallest unit for read and write (write is also called program, 1- \rightarrow 0)
- **Block** is the smallest unit for erase (0- \rightarrow 1)
 - i.e., make cells "empty" (i.e., no electrons)



Operation	Area
Read	Page
Program (Write)	Page
Erase	Block

Example



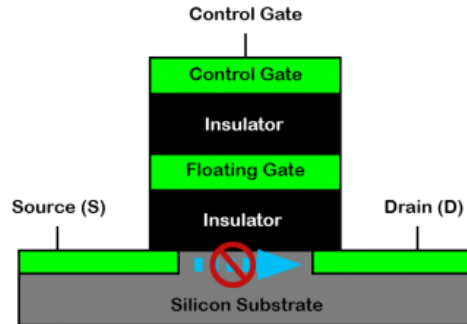
Latencies: read, write, and erase

	SLC	MLC	TLC	HDD	RAM
P/E cycles	100k	10k	5k	*	*
Bits per cell	1	2	3	*	*
Seek latency (μ s)	*	*	*	9000	*
Read latency (μ s)	25	50	100	2000-7000	0.04-0.1
Write latency (μ s)	250	900	1500	2000-7000	0.04-0.1
Erase latency (μ s)	1500	3000	5000	*	*
<i>Notes</i>	* metric is not applicable for that type of memory				
<i>Sources</i>	P/E cycles [20] SLC/MLC latencies [1] TLC latencies [23] Hard disk drive latencies [18, 19, 25] RAM latencies [30, 52] L1 and L2 cache latencies [52]				

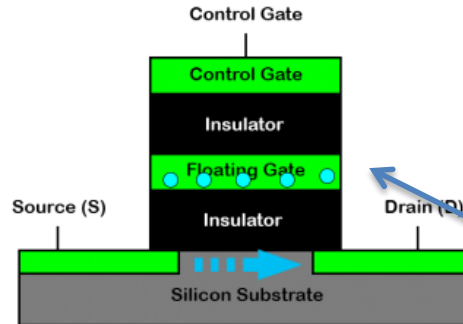
P/E cycle

- P: program/write; E: erase
- Every write & erase damages oxide layer surrounding the floating-gate to some extent
- P/E cycle:
 - Data are written to cells (P): cell value from 1 -> 0
 - Then erased (E): 0 -> 1


Multi-level cell (MLC)




State 1 - No Charge



State 2 - Lightly Charged



State 3 - Medium Charge



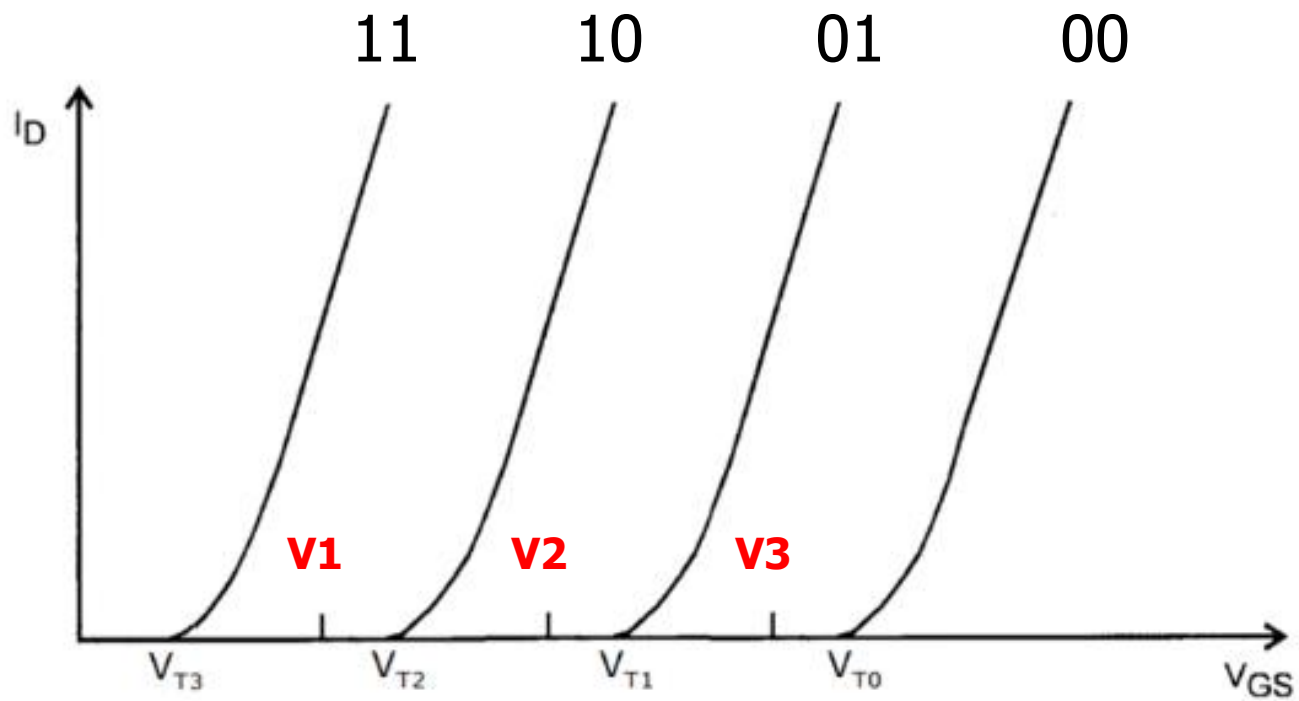
State 4 - Highly Charged

Note different levels of electrons

<http://www.cactus-tech.com/resources/blog/details/solid-state-drive-primer-2-slc-mlc-and-tlc-nand-flash>

MLC reading

- 2 bits, 3 intermediate voltages



Current vs voltage

SLC compared with MLC

- SLC:
 - Less complex
 - Faster
 - More reliable
 - Less storage
 - More costly

Read more

- [Solid-state revolution: in-depth on how SSDs really work](#)
- [How do SSDs work?](#)
 - <http://www.extremetech.com/extreme/210492-extremetech-explains-how-do-ssds-work>

References

- How Flash Memory Works
 - <https://www.youtube.com/watch?v=msi5GDz9JIw>
- Floating Gate Basics
 - <http://www.cse.scu.edu/~tschwarz/coen180/LN/flash.html>
- Friend of Flash
 - http://www.nnc3.com/mags/LM10/Magazine/Archive/2008/86/040-041_logfs/article.html

References

- Understanding Flash: Floating Gates and Wear
 - <https://flashdba.com/2015/01/09/understanding-flash-floating-gates-and-wear/>
- From Transistors to Functions
 - <http://www.cs.bu.edu/~best/courses/modules/Transistors2Gates/>

References

- Solid State Drive Primer
 - <https://www.cactus-tech.com/resources/blog/details/solid-state-drive-primer-1-the-basic-nand-flash-cell>
- How Does a Transistor Work?
 - <https://www.youtube.com/watch?v=lcrBqCFLHIY&feature=youtu.be>

References

- How Flash Memory Works
 - <https://www.youtube.com/watch?v=s7JLXs5es7I>